

# Document Clustering and Latent Semantic Indexing

CISC689/489-010, Lecture #18

Wednesday, April 22<sup>nd</sup>

Ben Carterette


## Clustering Review

- Cluster documents according to similarity in feature space
- Types of clustering:
  - Flat or hierarchical
  - “Hard” or “soft”
- Clustering algorithms:
  - K-means: flat, “hard” clusters
  - Agglomerative or divisive hierarchical: hierarchical, softer clusters

## K-Nearest Neighbor Clustering

- Alternative idea: fixed-size soft clusters
- *K-nearest neighbor*: for each item  $j$ , cluster it with the  $K$  things most similar to it
- K-nearest neighbor clustering forms one cluster per item
  - Clusters overlap
  - Does not necessarily have to cluster everything

## 5-Nearest Neighbor Clustering



## Evaluating Clustering

- Clustering will never be 100% accurate
  - Documents will be placed in clusters they don't belong in
  - Documents will be excluded from clusters they should be part of
  - A natural consequence of using term statistics to represent the information contained in documents
- Like retrieval and classification, clustering effectiveness must be evaluated
- Evaluating clustering is challenging, since it is an ***unsupervised*** learning task

## Evaluating Clustering

- If labels exist, can use standard IR metrics, such as precision and recall
  - In this case we are evaluating the ability of our algorithm to discover the “true” latent information

	Class A	Class B	Class C	Class D
Cluster 1	A <sub>1</sub>	B <sub>1</sub>	C <sub>1</sub>	D <sub>1</sub>
Cluster 2	A <sub>2</sub>	B <sub>2</sub>	C <sub>2</sub>	D <sub>2</sub>
Cluster 3	A <sub>3</sub>	B <sub>3</sub>	C <sub>3</sub>	D <sub>3</sub>
Cluster 4	A <sub>4</sub>	B <sub>4</sub>	C <sub>4</sub>	D <sub>4</sub>

$$prec_{cluster\ 1} = \frac{A_1}{A_1 + B_1 + C_1 + D_1}$$

$$rec_{cluster\ 1} = \frac{A_1}{A_1 + A_2 + A_3 + A_4}$$

- This only works if you have some way to “match” clusters to classes
- What if there are fewer or more clusters than classes?

## Evaluating Clusters

- “Purity”: the ratio between the number of documents from the dominant class in  $C$  to the size of  $C$

$$purity(C_i) = \frac{1}{|C_i|} \max_j |X \text{ s.t. } X \in C_i \text{ and } X \in K_j|$$

- $C_i$  is a cluster;  $K_j$  is a class
- Not such a great measure
  - Does not take into account coherence of the class
  - Optimized by making  $N$  clusters, one for each document

## Evaluating Clusters

- With no labeled data, evaluation is even more difficult
- Best approach:
  - Evaluate the system that the clustering is part of
  - E.g. if clustering is used to aid retrieval, evaluate the cluster-aided retrieval
- What kinds of systems use clustering?



## Clusters in IR Systems

- Automatic clustering has several uses:
  - Improving efficiency
  - Improving effectiveness of index
  - Improving effectiveness of retrieval

## Improving Efficiency

- Clusters can be a time- and space-saving device:
  - Cluster all documents in the corpus
  - Compare queries to cluster representations rather than document representations
  - Store cluster information in inverted lists rather than document information

$$t_i \rightarrow (cf_i, (c_1, tf_{1i}), \dots)$$

Cluster frequency   Cluster term frequency

- Important in the 70s and 80s, not so much today

## Improving Effectiveness

- Recall the cluster hypothesis:
  - “Closely associated documents tend to be relevant to the same requests”
- We may be able to improve retrieval by finding documents that are closely associated with documents that are likely to be relevant
  - Even if they don’t contain query terms
  - E.g. perhaps they contain related terms the user didn’t think of (“industry” → “companies”, “businesses”, “producers”, ...)

## Improving Effectiveness by Ranking Clusters


- As with clustering for efficiency, cluster all documents before indexing
- Store cluster information in inverted lists (but keep document information too)
- When a user enters a query, score the clusters
- Then score documents within the top-scoring clusters

## Improving Effectiveness by Ranking Clusters

- Two approaches:
  - Rank the clusters from highest scoring to lowest scoring; within each cluster rank documents from highest scoring to lowest scoring
  - Rank the documents in the K highest-scoring clusters from highest score to lowest score
- Both tend to find relevant documents that are not found with non-clustering methods
  - But also miss relevant documents that are found with non-clustering methods

## Scoring Clusters

- A cluster can be scored the same way as a document
  - Vector space model: cosine similarity between query vector and cluster centroid
  - Language model:  $P(Q|C) = \prod_{t \in Q} P(t|C) = \prod_{t \in Q} (1 - \alpha_C) \frac{tf_{tC}}{|C|} + \alpha_C \frac{ctf_t}{|G|}$ 



Smoothing with collection
- Or other ways:
  - $S(C, Q) = \max_{D_i \in C} S(D_i, Q)$
  - $S(C, Q) = \min_{D_i \in C} S(D_i, Q)$
  - $S(C, Q) = \frac{1}{|C|} \sum_{D_i \in C} S(D_i, Q)$

## Using Clusters to Adjust Document

- Language models require background to smooth with
- Clusters provide a background, perhaps more focused than using entire collection
- Smooth document language model with cluster background first, then with collection background
  - $P(w | D)$  – frequency of term in document
  - $P(w | C)$  – frequency of term in all documents in a cluster
  - $P(w | G)$  – frequency of term in all documents in the collection

## Smoothing Document Scores

- Basic approach: linear interpolation

$$P(w|D) = \lambda_1 \frac{tf}{|D|} + \lambda_2 \frac{\sum_{D \in C} tf}{\sum_{D \in C} |D|} + \lambda_3 \frac{ctf}{|G|}$$

- Better approach: smooth cluster with collection, then smooth document with both

$$P(w|D) = \lambda \frac{tf}{|D|} + (1 - \lambda) \left( \beta \frac{\sum_{D \in C} tf}{\sum_{D \in C} |D|} + (1 - \beta) \frac{ctf}{|G|} \right)$$



## Query-Time Clustering

- Clustering the entire collection takes a long time
- Can only use simple algorithms like K-means
- Instead, cluster the top documents ranked for a query
- Use those clusters to re-score and re-rank the documents

## Does it Work?

- Retrieving clusters:

Verdict: no apparent improvement

Collection	First-stage doc retrieval (QL+DM)	Group-average	Single-linkage	Complete-linkage	Centroid
AP (training)	0.2179	0.2161 (t=0.8)	0.2153 (t=0.8)	0.2130 (t=0.8)	0.2164 (t=0.7)
WSJ	0.2958	0.2902 (t=0.8)	0.2911 (t=0.8)	0.2889 (t=0.8)	0.2936 (t=0.8)

- Clustering at index time, use clusters for smoothing

Verdict: improvement over baselines

Collection	Simple Okapi	QL+DM	QL+CBDM	%chg
AP (K=2000)	0.2198	0.2179	0.2326 (+)	+6.73*
WSJ (K=2000)	0.2762	0.2958 (+)	0.3006 (+)	+1.62*
FT (K=2000)	0.2556	0.2610	0.2713 (+)	+3.95*
SJMN (K=2000)	0.2098	0.2032	0.2171 (+)	+6.88*
LA (K=2000)	0.2279	0.2468 (+)	0.2590 (+)	+4.94*
FR (K=1000)	0.2644	0.2875	0.3316	+15.37

## Does it Work?

- Clustering at query time, use clusters for smoothing

Collection	Threshold	QL+TDM	QL+CBDM	%chg
AP	0.2	0.2107	0.2223	+5.46*
	0.4	0.2140	0.2247	+5.02*
	0.6	0.2113	0.2211	+4.64*
WSJ	0.2	0.2663	0.2954	+10.92*
	0.4	0.2707	0.3004	+10.95*
	0.6	0.2685	0.2998	+11.65*
FR	0.2	0.2409	0.2935	+21.84
	0.4	0.2265	0.2710	+19.64
	0.6	0.2276	0.2933	+28.84

Verdict: improvement  
over baselines

## Clusters in IR – Summary

- Index-time clustering
  - Saves space in inverted file
  - Build topical hierarchies
  - Retrieve clusters of documents rather than individual documents
- Query-time clustering
  - After query is submitted, cluster results
  - Possibly detect subtopics or different interpretations of query

## Latent Semantic Indexing

- Clustering only puts documents together based on term similarity
- Can we do more than that?
  - We'd like synonyms and highly related terms to count for more when calculating document similarity
  - And words that have multiple senses to count for less
- How can we do this?

## Linear Algebra Background

- A *vector space* is defined by a set of linearly independent *basis vectors*
  - Linearly independent: no vector can be expressed as a linear combination of other vectors
- Every vector can be expressed as a linear combination of the basis vectors

- **Example:**

These three vectors form a  
3-dimensional vector space

$$B = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

These three vectors form a  
2-dimensional vector space

$$B = \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 2 \\ 2 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

## Bases in the Vector Space Model

- When we discussed the VSM, we assumed bases could be formed from terms
  - Document and query vectors are linear combinations of term basis vectors
- But basis vectors have to be linearly independent—do terms satisfy this?
  - Probably not
  - Two terms that always appear together are not independent
  - More insidious example:
    - “bush” appears in documents about landscaping and documents about politics
    - The vector for “bush” may be a linear combination of many vectors for terms related to landscaping and terms related to politics

## Bases in the Vector Space Model

- Is there a better way to choose the bases?
- *Semantic concepts*
  - Find a group of topics or concepts that are orthogonal
    - E.g. “landscaping” and “politics” topics are probably linearly independent
  - Each concept forms a basis vector
  - A document vector is a linear combination of concept vectors
- Note: this is not easy to do
  - Solving this problem would probably solve all of AI

## Finding Linearly Independent Bases

- There are methods for finding linearly independent basis vectors
- We will apply one method and assume that the bases it produces represent “concepts”

## Linear Algebra Background

- Matrix-vector multiplication:

$$Ax = \begin{bmatrix} a_{11} & a_{12} & \dots \\ a_{21} & a_{22} & \dots \\ \dots & \dots & \dots \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \dots \end{bmatrix} = \begin{bmatrix} a_{11}x_1 + a_{12}x_2 + \dots \\ a_{21}x_1 + a_{22}x_2 + \dots \\ \dots \end{bmatrix} = x_1 \begin{bmatrix} a_{11} \\ a_{21} \\ \dots \end{bmatrix} + x_2 \begin{bmatrix} a_{12} \\ a_{22} \\ \dots \end{bmatrix}$$

–  $A$  transforms  $x$

- *Eigenvalues and eigenvectors*

– Given a square matrix  $A$ , the eigenvectors of  $A$  are the vectors  $x$  such that  $Ax$  is a scalar multiple of  $x$

- i.e. the vectors that are only transformed by length, not by direction

$$Ax = \lambda x$$

Every  $x$  that satisfies this equation is an eigenvector.  
The *eigenvalues*  $\lambda$  show how much  $A$  shortens or elongates  $x$ .

## Eigenvectors and Eigenvalues

- Example:

$$\begin{bmatrix} 6 & -2 \\ 4 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 2 \end{bmatrix} = \begin{bmatrix} 2 \\ 4 \end{bmatrix} = 2 \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

- $\begin{bmatrix} 1 & 2 \end{bmatrix}'$  is an eigenvector of the matrix, and 2 is an eigenvalue
- How many eigenvalues are there for an  $n$  by  $n$  matrix?

$$Ax = \lambda x \Leftrightarrow (A - \lambda I)x = 0$$

- $x$  is nonzero only if *determinant* of  $A - \lambda I = 0$
- The determinant is a polynomial in  $\lambda$  of degree  $n$
- Therefore there are at most  $n$  eigenvalues

## Examples

- Example 1:

$$A = \begin{bmatrix} 30 & 0 & 0 \\ 0 & 20 & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad \begin{array}{l} \text{eigenvalues} = 30, 20, 1 \\ \text{eigenvectors} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} \end{array}$$

- Example 2:

$$A = \begin{bmatrix} 0.71 & 0.03 & -0.01 \\ 0.03 & 1.02 & 0.03 \\ -0.01 & 0.03 & 0.70 \end{bmatrix} \quad \begin{array}{l} \text{eigenvalues} = 1.03, 0.72, 0.69 \\ \text{eigenvectors} = \begin{bmatrix} -0.09 \\ -0.99 \\ -0.09 \end{bmatrix}, \begin{bmatrix} 0.82 \\ -0.03 \\ -0.57 \end{bmatrix}, \begin{bmatrix} 0.56 \\ -0.12 \\ 0.82 \end{bmatrix} \end{array}$$

- Exercise: verify that  $Ax = \lambda x$  holds

## Eigenvectors and Eigenvalues

- Eigenvectors of symmetric matrices are *linearly independent*
  - Why? If an eigenvector  $x$  could be written as a sum of other vectors, then  $Ax = \lambda x$  would not be true— $x$  would not be an eigenvector in the first place!
- Therefore eigenvectors of symmetric matrices form a basis
  - Every vector in the space is a linear combination of the eigenvectors

## Eigenvectors and Eigenvalues

- Eigenvalues of real-valued matrices are real numbers
- Eigenvalues of *positive semidefinite* matrices are non-negative

## Eigen Decompositions

- A square matrix  $A$  can be decomposed into a matrix product  $A = U\Lambda U^{-1}$  where
  - $U$  is a matrix with eigenvectors of  $A$  as columns
  - $\Lambda$  is a matrix with eigenvalues of  $A$  in decreasing order on the diagonal and 0 elsewhere
- A symmetric square matrix  $A$  can be decomposed into a matrix product  $A = Q\Lambda Q'$ 
  - $Q$  is a real orthogonal matrix with normalized eigenvectors as columns

## So How Does This Apply to IR?

- If we had the right kind of matrix, a decomposition might be able to find orthogonal “concepts” in the document/term data
- Those concepts might be a better basis for a vector space model
- We could take the  $K$  most important concepts (corresponding to the  $K$  greatest eigenvalues) to reduce the dimensionality of the space
- ... but we don't have square symmetric matrices in IR



## Eigen Decomp in IR

- Arrange  $N$  documents and  $V$  terms into a  $V \times N$  matrix called  $A$ , where  $A_{ij}$  = term weight of term  $i$  in document  $j$ 
  - This is neither square nor symmetric
- We can compute the following matrix products:
  - $AA'$  = a  $V \times V$  matrix of document similarities
  - $A'A$  = a  $N \times N$  matrix of term similarities
  - Note that these are both square and symmetric—both have eigen decompositions


## Singular Value Decomposition

- SVD uses  $AA'$  and  $A'A$  to compute an eigen decomposition of  $A$ 
  - $AA'$  and  $A'A$  have the same number  $m$  of eigenvectors
    - $m \leq \min(N, V)$
- Specifically:  $A = U\Sigma V'$ 
  - $U$  = a  $V \times m$  matrix with eigenvectors of  $AA'$  as cols
  - $V$  = a  $N \times m$  matrix with eigenvectors of  $A'A$  as cols
  - $\Sigma$  = an  $m \times m$  matrix with the square roots of eigenvalues of  $AA'$  on the diagonal
    - Eigenvalues of  $AA'$  = eigenvalues of  $A'A$

## SVD on the Document-Term Matrix

- Eigenvectors of  $AA'$  represent “document concepts”
- Eigenvectors of  $A'A$  represent “term concepts”
- Eigenvalues give the relative weight of the concepts
- The occurrence of a term in a document is a linear combination of “term concepts” and “document concepts”

### Illustration



Each column of  $U$  is a “document concept”

Each row of  $V'$  is a “term concept”

$$A_{ij} = \sum_{k=1}^m U_{ik} V_{kj}$$

## Example

### Technical Memo Example

**Titles:**

- c1: *Human machine interface for Lab ABC computer applications*
- c2: *A survey of user opinion of computer system response time*
- c3: *The EPS user interface management system*
- c4: *System and human system engineering testing of EPS*
- c5: *Relation of user-perceived response time to error measurement*

m1: The generation of random, binary, unordered *trees*  
 m2: The intersection *graph* of paths in *trees*  
 m3: *Graph minors* IV: Widths of *trees* and well-quasi-ordering  
 m4: *Graph minors*: A survey

Terms	Documents							
	c1	c2	c3	c4	c5	m1	m2	m4
human	1	0	0	1	0	0	0	0
interface	1	0	1	0	0	0	0	0
computer	1	1	0	0	0	0	0	0
user	0	1	1	0	1	0	0	0
system	0	1	1	2	0	0	0	0
response	0	1	0	1	0	0	0	0
time	0	1	0	0	1	0	0	0
EPS	0	0	1	1	0	0	0	0
survey	0	1	0	0	0	0	0	0
trees	1	0	0	0	1	1	1	0
graph	0	0	0	0	0	1	1	1
minors	0	0	0	0	0	0	1	1

From Deerwester, Dumais, Harshman, "Indexing by Latent Semantic Analysis"

## Example

Figure 1 displays two heatmaps showing the correlation matrix of the 18 variables. The top heatmap is for  $T_0$  and the bottom for  $D_0$ . Both show a 18x18 matrix of correlation coefficients with a color scale from -0.5 (blue) to 0.5 (red). The variables are ordered as follows:  $S_0$ ,  $3.34$ ,  $2.54$ ,  $2.35$ ,  $1.64$ ,  $1.50$ ,  $1.31$ ,  $0.85$ ,  $0.56$ ,  $0.36$ ,  $0.27$ ,  $0.27$ ,  $0.30$ ,  $0.21$ ,  $0.01$ ,  $0.04$ ,  $0.03$ ,  $0.20$ ,  $0.61$ ,  $0.54$ ,  $0.28$ ,  $0.00$ ,  $0.01$ ,  $0.02$ ,  $0.01$ ,  $0.02$ ,  $0.08$ ,  $0.53$ .

Verify:  $T_0 S_0 D_0 =$  original document-term matrix

## Optimal Dimensionality Reduction

- SVD can also be used for dimensionality reduction
- Reduce  $\Sigma$  to the top-k largest eigenvalues
- For documents and terms, this effectively reduces the dimensionality to the k most important “concepts”
  - Furthermore, the reduction is optimal—it is the best reduction you could possibly do given the document-term matrix

## Reducing to k=2 Concepts

$X \approx$	$T$	$S$	$D'$
0.22	-0.11	3.34	0.20 0.61 0.46 0.54 0.28 0.00 0.02 0.02 0.08
0.20	-0.07	2.54	-0.06 0.17 -0.13 -0.23 0.11 0.19 0.44 0.62 0.53
0.24	0.04		
0.40	0.06		
0.64	-0.17		
0.27	0.11		
0.27	0.11		
0.30	-0.14		
0.21	0.27		
0.01	0.49		
0.04	0.62		
0.03	0.45		
$X_{\text{ihat}} =$			
0.16	0.40	0.38	0.47 0.18 -0.05 -0.12 -0.16 -0.09
0.14	0.37	0.33	0.40 0.16 -0.03 -0.07 -0.10 -0.04
0.15	0.51	0.36	0.41 0.24 0.02 0.06 0.09 0.12
0.26	0.84	0.61	0.70 0.39 0.03 0.08 0.12 0.19
0.45	1.23	1.05	1.27 0.56 -0.07 -0.15 -0.21 -0.05
0.16	0.58	0.38	0.42 0.28 0.06 0.13 0.19 0.22
0.16	0.58	0.38	0.42 0.28 0.06 0.13 0.19 0.22
0.22	0.55	0.51	0.63 0.24 -0.07 -0.14 -0.20 -0.11
0.10	0.53	0.23	0.21 0.27 0.14 0.31 0.44 0.42
-0.06	0.23	-0.14	-0.27 0.14 0.24 0.55 0.77 0.66
-0.06	0.34	-0.15	-0.30 0.20 0.31 0.69 0.98 0.85
-0.04	0.25	-0.10	-0.21 0.15 0.22 0.50 0.71 0.62

## Latent Semantic Analysis for Retrieval

$X \approx$	$T$	$S$	$D'$
0.22	-0.11	3.34	0.20 0.61 0.46 0.54 0.28 0.00 0.02 0.02 0.08
0.20	-0.07	2.54	-0.06 0.17 -0.13 -0.23 0.11 0.19 0.44 0.62 0.53
0.24	0.04		
0.40	0.06		
0.64	-0.17		
0.27	0.11		
0.27	0.11		
0.30	-0.14		
0.21	0.27		
0.01	0.49		
0.04	0.62		
0.03	0.45		

- Now  $D$  contains the new document vectors
  - Each a 2-D vector of “concept weights”
- To process a query  $Q$ :
  - Use  $T$  to transform it to the 2-D concept space
  - Weight the concepts using  $S$
  - $Q' = Q' T S^{-1}$
  - Calculate cosine similarity between  $Q'$  and each document

## Example

$X \approx$	$T$	$S$	$D'$
0.22	-0.11	3.34	0.20 0.61 0.46 0.54 0.28 0.00 0.02 0.02 0.08
0.20	-0.07	2.54	-0.06 0.17 -0.13 -0.23 0.11 0.19 0.44 0.62 0.53
0.24	0.04		
0.40	0.06		
0.64	-0.17		
0.27	0.11		
0.27	0.11		
0.30	-0.14		
0.21	0.27		
0.01	0.49		
0.04	0.62		
0.03	0.45		

$Q$  = “human computer”

$$Q' = [ 1 \ 0 \ 1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 ]$$

$$Q'T = [ 1*0.22+1*0.24 \quad 1*-0.11+1*0.04 ]$$

$$= [ 0.46 \quad -0.07 ]$$

$$Q'TS^{-1} = [ 0.46/3.34 \quad -0.07/2.54 ] = [ 0.14 \quad -0.03 ]$$

## LSI: Does it Work?

- Seems to be a little better than standard vector space model
  - Recall is good: intuitively it is retrieving “clusters” of documents related by topic, which improves recall
  - Precision is OK: it can find things that are not really related
- When it does not do well, it is hard to understand what happened
- It takes a very long time to do the SVD
  - In 1990, one day for ~10,000 documents