

PageRank

CISC489/689-010, Lecture #20

Wednesday, April 29th

Ben Carterette

Web Search

- Problem:
 - Web search engines easily hacked
 - I want to sell something; I'll just add a few popular keywords to my page over and over and over again
 - All the retrieval models we've discussed will score that page higher for those keywords
- Other problems:
 - No hackers, but top-ranked pages are coming from deep within a site, or from pages that change often, or pages about very obscure topics
 - Not really useful

Possible Solution

- Leverage link structure
- Maybe if many pages are linking to a page, that page is more “important”
- Idea:
 - Count the number of inlinks to the page
 - Assign it “importance” based on that number
- Any problem with this?


Hacking Link Counts

- I can just make a bunch of pages that link to my spam page
- Inlink count will be high even though my page is not important
- Better idea:
 - Recursively use the importance of the linking pages when calculating the importance of the page

PageRank

- Google's PageRank is probably the best known algorithm
- Intuitive idea: "random surfer" model
 - If you start on a random page on the internet and just start clicking links randomly,
 - What is the probability you will land on page u ?
 - If one page has a higher landing probability, the pages it links to have higher landing probabilities as well
 - Higher probability = more authority = better PageRank

Illustration



PageRank Definition

$$R(u) = c \sum_{v \in B_u} \frac{R(v)}{N_v}$$

$R(u)$ is the PageRank of u

B_u is a set of pages that link to u

v is one of the pages in B_u

$R(v)$ is the PageRank of v

N_v is the number of pages that v links to

Sinks

- Problem:
 - I have two pages that only link to each other, plus one page that links to one of them
 - When the “random surfer” gets to one of those pages, he will just keep alternating between them
 - Their PageRank will dominate everything else
- Solution: once in a while the random surfer just starts over at a new page
 - OK, but how do I put that in PageRank?

Random Restarts

$$R'(u) = c \sum_{v \in B_u} \frac{R'(v)}{N_v} + cE(u)$$

$E(u)$ is the probability that a random surfer jumps to u

Calculating PageRank

- A simple iterative algorithm:
 - First, assign a PageRank to every page
 - E.g. $R_0(u) = 1/N$
 - Initialize $E(u) = \alpha$ for all u ($0.15/N$ in original paper)
 - Over iterations $i=1\dots$, do
 - Update each PageRank as: $R_{i+1}(u) = \sum_{v \in B_u} \frac{R_i(v)}{N_v}$
 - Calculate d as the sum of PageRanks from the previous iteration minus the sum of PageRanks from the current iteration
 - Update each PageRank as $R_{i+1}(u) = R_{i+1} + d\alpha$
 - Calculate δ as the sum of |PageRanks from the current iteration minus PageRanks from the previous iteration|
 - If $\delta > \epsilon$, PageRanks have converged

Scalability

- Calculating PageRank requires a vector for every URL
- Along with a list of the URLs that link to that URL and a list of URLs that page links to
- Space usage is $O(N^2)$
- Time complexity also $O(N^2)$
- Even for small collections (like Wikipedia), it is nearly impossible to keep all of this in memory

Calculating PageRank with MapReduce

- First: extract links
 - For every page, I need to know all the pages that link to it
 - MapReduce solution:
 - Map operator takes page u and outputs (v, u) for every URL v in page u
 - Reduce operator takes all tuples with key v and reduces them to a list of unique URLs B_u that link to v
 - $(v, (u_1, u_2, u_3, \dots))$

Calculating PageRank with MapReduce

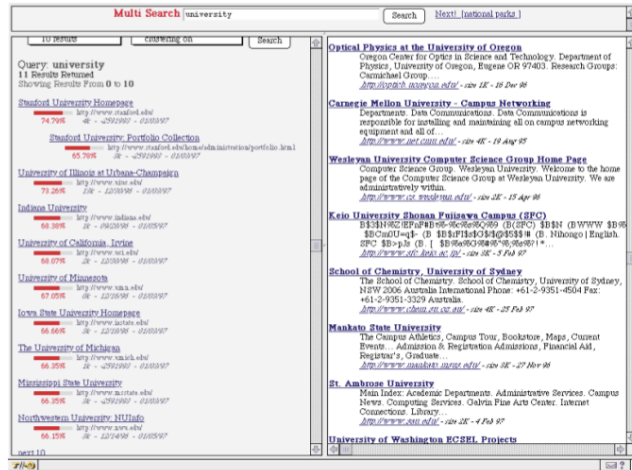
- Next: calculate PageRank iteratively
 - First, initialize all PageRanks to $1/N$
 - Then iteratively MapReduce
 - Map operator takes a page u_i and the URLs v_j ($j=1..n$) that it links to, and outputs $(v_j, R(u_i)/n)$
 - Reduce operator takes pairs $(v_j, R(u_i)/n)$ and outputs

$$\left(v_j, (1-d) + d \sum_{i=1}^m \frac{R(u_i)}{n} \right)$$
 - Calculate delta and determine whether converged
 - If not, MapReduce again

PageRank for IR

- PageRank is query-independent
 - The importance of a page is not related to any query
 - We cannot simply rank pages by PageRank
- PageRank can be used to re-rank results that have been retrieved for a query
- It can also be used as a feature in the ranking function
- Or as a weight on anchor text features

Example Use of PageRank



From “The PageRank Citation Algorithm: Bringing Order to the Web”, Page et al.

Wikipedia PageRanks

Page Title	PageRank
United States	2.9×10^{-3}
France	1.3×10^{-3}
United Kingdom	1.2×10^{-3}
England	1.0×10^{-3}
Germany	1.0×10^{-3}
Canada	0.9×10^{-3}
2007	0.8×10^{-3}
World War II	0.8×10^{-3}
Australia	0.7×10^{-3}
2008	0.7×10^{-3}

Based on links between Wikipedia pages

A Bit of Theory

- Markov chain:
 - N states with transition probability matrix P
 - At any time we are in exactly one state
 - P_{ij} indicates probability of moving from state i to state j
 - For all i, $\sum_{j=1}^n P_{ij} = 1$

A Bit of Theory

- Ergodic Markov chains
 - *Ergodic* means there is a path between any two states
 - No matter what state you start in, the probability of being in any other state after T steps is greater than zero (after a *burn-in* time T_0)
 - Over many steps T, each state has some “visit rate”: starting from any state, we will visit each state according to its visit rate
 - This is the *steady state* for the chain

A Bit of Theory

- Let x_0 be a vector representing our current state

$$x_0 = [0 \quad 0 \quad 0 \quad \dots \quad 1 \quad \dots \quad 0 \quad 0]$$

1 in position i , 0s everywhere else

- What is the probability of each possible state we can transition to from x_0 ?

$$x_1 = x_0 P$$

- And the probability of each possible state from x_1 (two steps from x_0)?

$$x_2 = x_1 P = (x_0 P) P = x_0 P^2$$

A Bit of Theory

- After k steps, $x_k = x_0 P^k$
- As k goes to infinity, x_k converges to the steady state
- When x_k is the steady state, $x_k P = x_k$
- The steady state is an *eigenvector* of P
 - As it turns out, it is the *principle eigenvector*
 - Eigenvalue = 1
- If P is a matrix of links between documents, the principle eigenvector holds the PageRanks

PageRank Modifications

- The $E(u)$ quantities solve the sink problem, but can also be used to adjust PageRanks
- Usually, $E(u)$ assigned uniformly
 - Equal probability to jump to any page
- Instead, bias to certain pages
 - One possibility: assign one page $E(u) = \alpha$, all other pages 0
 - For example, Yahoo home page
 - Then when the “random surfer” restarts, she always restarts at the same place
 - Result: Yahoo gets highest PageRank, followed by pages Yahoo links to

Topic-Based PageRank

- A different kind of random surfer:
 - First picks a category randomly
 - Then jumps to a page randomly within that category
- Instead of calculating a single PageRank for each page, calculate M PageRanks, one for each category
 - Category PageRank = PageRank among other pages in the same category

Topic-Based PageRank for Personalization

- Each individual user is more interested in some categories than others
- Calculate the probability that a user is interested in a category based on the frequency they visit pages in that category
 - E.g. sports = 0.7, finance = 0.2, health = 0.1, all others = 0
- Then the personalized PageRank for a page u is the weighted sum of category PageRanks for that page

$$R(u) = \sum_{i=1}^M p_i R_i(u)$$

p_i is the user's category i probability
 $R_i(u)$ is the PageRank of u for category i

Hyperlink-Induced Topic Search (HITS)


- Another link-graph algorithm
- Idea:
 - Some pages are *authoritative*: they are very informative about a topic
 - Other pages are *hubs* for a topic: they link to a lot of pages on the topic
- Example:
 - CiteSeer links to a lot of computer science research papers —it's a computer science research hub
 - The papers it links to are computer science research authorities
- Find both hubs and authorities

Hubs and Authorities

- Hubs are pages that link to a lot of authorities
- Authorities are pages that are linked to by a lot of hubs
- Another recursive definition

HITS Algorithm

- First, get a *root set* of pages
 - Pages that match a query, for example
- From the root set, construct a *base set*
 - Pages that link to the root set and pages that the root set link to



From "Authoritative Sources in a Hyperlinked Environment", J. Kleinberg

Figure 1: Expanding the root set into a base set.

HITS Algorithm

- Initialize “hub score” $h(u)=1$ and “authority score” $a(u)=1$ for each page u in the base set
- Then iteratively update $h(u)$ and $a(u)$ for all u :

$$h(u) = \sum_{v \in F_u} a(v)$$

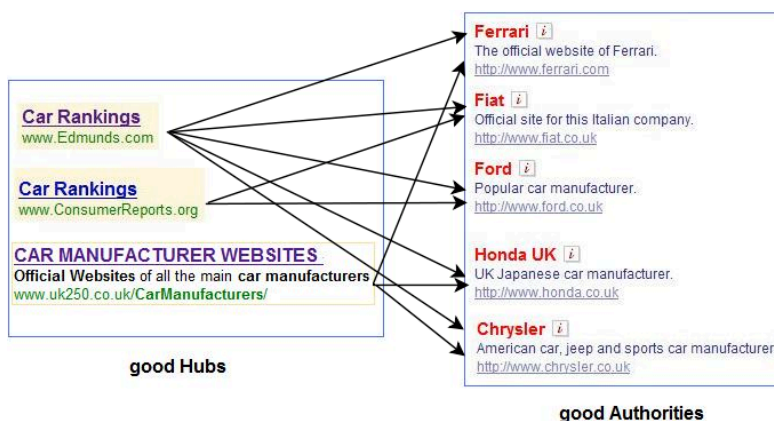
F_u = set of pages that u links to

$$a(u) = \sum_{v \in B_u} h(v)$$

B_u = set of pages that link to u

- After each iteration, divide $h(u)$ and $a(u)$ by some constant
- After only a few iterations, scores converge

HITS Example Results



Query: **Top automobile makers**

From <http://www.math.cornell.edu/~mec/Winter2009/RalucaRemus/Lecture4/Images/cars1.png>

Matrix Form

- P is the transition matrix
- PP' is a sort of “similarity” matrix in terms of links to other pages
 - Entry i,j is higher if pages i and j link to the same pages
- $P'P$ is a sort of “similarity” matrix in terms of links from other pages
 - Entry i,j is higher if pages i and j are linked to from the same pages

Matrix Form

- We can write the hub score as a matrix-vector product:
 - $h = Pa$ (transition matrix times authority score)
- We can write authority score as
 - $a = P'h$ (transpose of the transition matrix times hub score)
- Substituting, we get
 - $h = PP'h$
 - $a = P'Pa$

Matrix Eigenvectors

- If $h = PP'h$, then h is an eigenvector of the “outlink similarity matrix”
 - Hub scores are a basis vector of a space defined by outlinks
- If $a = P'Pa$, then a is an eigenvector of the “inlink similarity matrix”
 - Authority scores are a basis vector of a space defined by inlinks