

CS3245

# Information Retrieval

Lecture 6: Index Compression

6

# Last Time: index construction

---

- Sort-based indexing
  - Blocked Sort-Based Indexing
    - Merge sort is effective for disk-based sorting (avoid seeks!)
  - Single-Pass In-Memory Indexing
    - No global dictionary - Generate separate dictionary for each block
    - Don't sort postings - Accumulate postings as they occur
- Distributed indexing using MapReduce
- Dynamic indexing: Multiple indices, logarithmic merge



# Today: Cmprssn

BRUTUS	→	1	2	4	11	31	45	173	174
--------	---	---	---	---	----	----	----	-----	-----

CAESAR	→	1	2	4	5	6	16	57	132	...
--------	---	---	---	---	---	---	----	----	-----	-----

CALPURNIA	→	2	31	54	101
-----------	---	---	----	----	-----

- Collection statistics in more detail (with RCV1)
  - How big will the dictionary and postings be?
- Dictionary compression
- Postings compression



# Why compression (in general)?

- Use less disk space
  - Saves a little money
- Keep more data in memory
  - Increases speed
- Increase speed of data transfer from disk to memory
  - [read compressed data | decompress] is faster than [read uncompressed data]
  - Premise: Decompression algorithms are fast
    - True of the decompression algorithms we use



# Lossless vs. lossy compression

- Lossless compression: All information is preserved
  - What we mostly do in IR.
- Lossy compression: Discard some information
- Several of the preprocessing steps can be viewed as lossy compression: case folding, stop words, stemming, number elimination
- Later: Prune postings entries that are unlikely to turn up in the top  $k$  list for any query
  - Almost no loss quality for top  $k$  list



# Vocabulary vs. collection size

- Heaps' law:  $M = kT^b$
- $M$  is the size of the vocabulary,  $T$  is the number of tokens in the collection
- Typical values:  $30 \leq k \leq 100$  and  $b \approx 0.5$
- In a log-log plot of vocabulary size  $M$  vs.  $T$ , Heaps' law predicts a line with slope about  $\frac{1}{2}$ 
  - It is the simplest possible relationship between the two in log-log space
  - An empirical finding (“empirical law”)

# Heaps' Law

For RCV1, the dashed line


$$\log_{10} M = 0.49 \log_{10} T + 1.64$$

is the best least squares fit.

Thus,  $M = 10^{1.64} T^{0.49}$  so  $k = 10^{1.64} \approx 44$  and  $b = 0.49$ .

Good empirical fit for  
Reuters RCV1 !

For first 1,000,020 tokens,  
law predicts 38,323 terms;  
actually, 38,365 terms



# Zipf's law



- How about the relative frequencies of terms?
- In natural language, there are a few very frequent terms and very many very rare terms.
- Zipf's law: The  $i$ th most frequent term has frequency proportional to  $1/i$ .
- $cf_i \propto 1/i = K/i$  where  $K$  is a normalizing constant
- $cf_i$  is collection frequency (not document frequency): the number of occurrences of the term  $t_i$  in the collection.






# Zipf consequences

- If the most frequent term (*the*) occurs  $cf_1$  times
  - then the second most frequent term (*of*) occurs  $cf_1/2$  times
  - the third most frequent term (*and*) occurs  $cf_1/3$  times ...
- Equivalent:  $cf_i = K/i$  where  $K$  is a normalizing factor, so  $\log cf_i = \log K - \log i$ 
  - Linear relationship between  $\log cf_i$  and  $\log i$
- Another power law relationship

# Zipf's law for Reuters RCV1





# DICTIONARY COMPRESSION



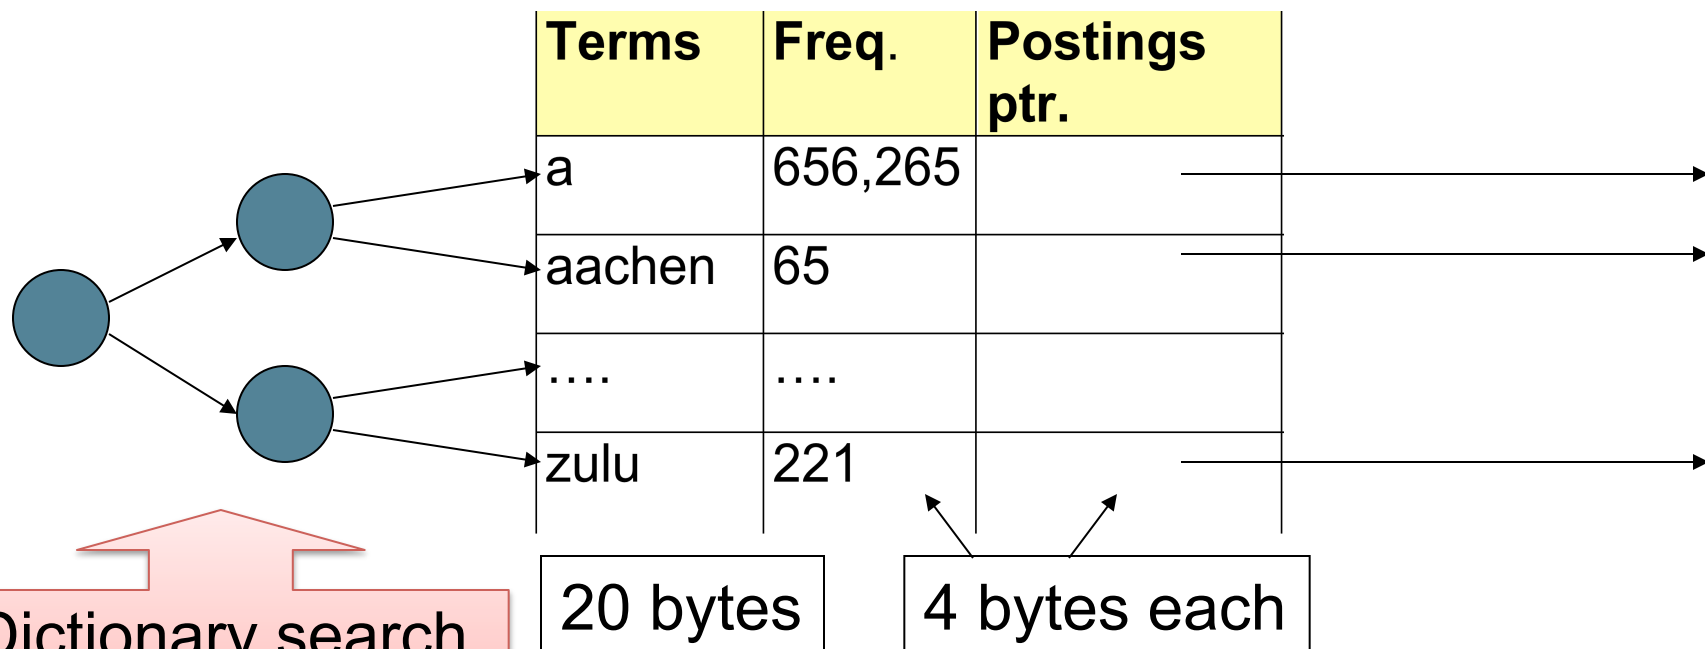
# Why compress the dictionary?

- Search begins with the dictionary
- We want to keep it in memory
- Memory footprint competition with other applications
- Embedded/mobile devices may have very little memory
- Even if the dictionary isn't in memory, we want it to be small for a fast search startup time

***Compressing the dictionary is important***

# Dictionary storage - first cut

- Array of fixed-width entries
  - ~400,000 terms; 28 bytes/term = 11.2 MB.





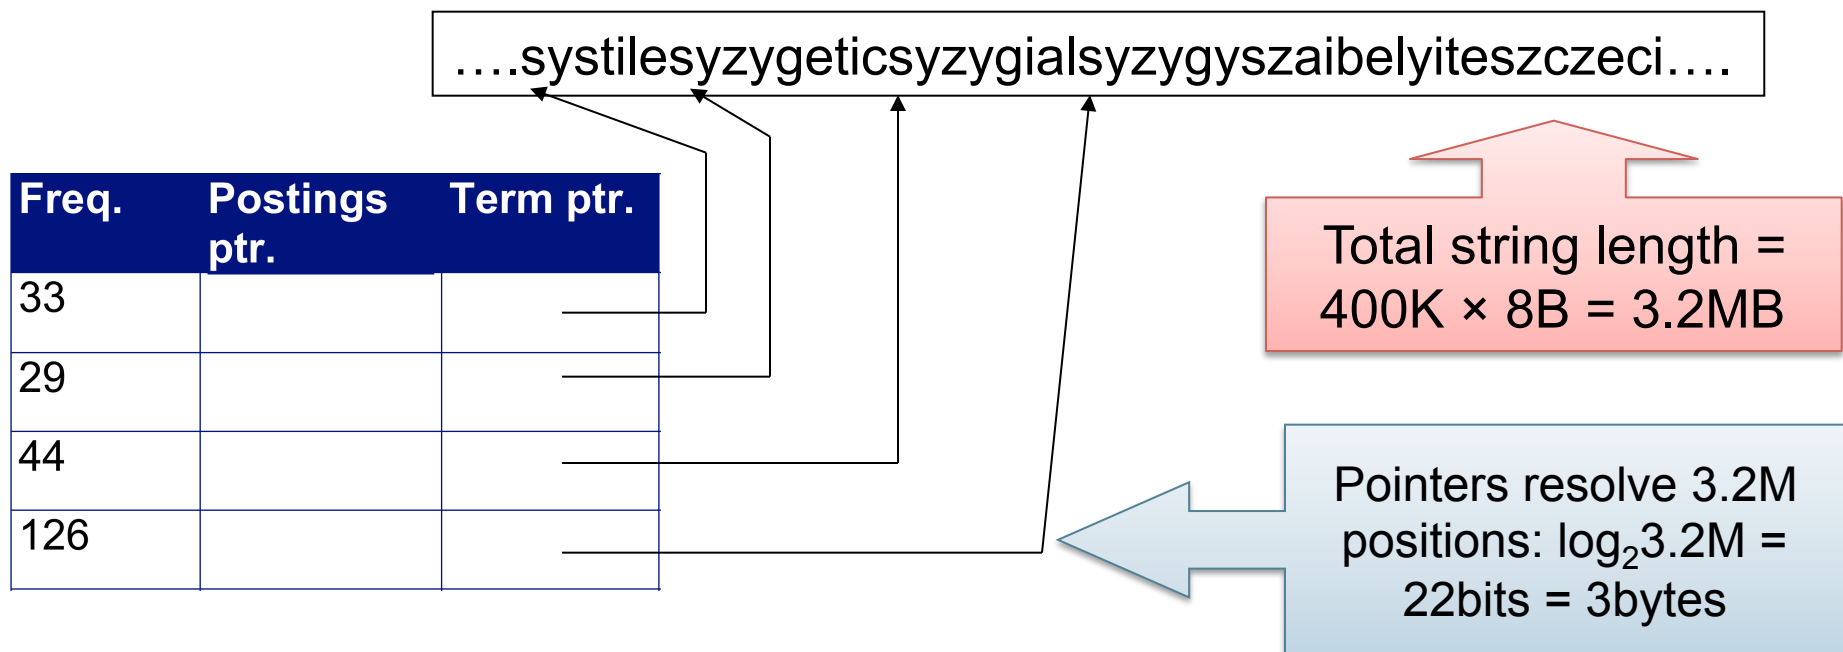
# Fixed-width terms are wasteful

- Most of the bytes in the **Term** column are wasted – we allot 20 bytes for 1 letter terms.
  - And we still can't handle *supercalifragilisticexpialidocious* or *hydrochlorofluorocarbons*.
- Written English averages ~4.5 characters/word.
- Average dictionary word in English: ~8 characters
  - How do we use ~8 characters per dictionary term?
- Short words dominate token counts but not type average.

# Compressing the term list: Dictionary-as-a-String



- Store dictionary as a (long) string of characters:
  - Pointer to next word shows end of current word
  - Hope to save up to 60% of dictionary space.





# Space for dictionary as a string

- 4 bytes per term for frequency
  - 4 bytes per term for pointer to postings
  - 3 bytes per term pointer
  - Avg. 8 bytes per term in term string
  - $400\text{K terms} \times 19 \Rightarrow 7.6 \text{ MB}$  (against 11.2MB for fixed width)
- } Now avg. 11 bytes/term, not 20.





# Blocking

- Store pointers to every  $k$ th term string.
  - Example below:  $k=4$ .
- Need to store term lengths (1 extra byte)

....**7***systile***9***syzygetic***8***syzygial***6***syzygy***11***szaibelyite* ...

Freq.	Postings ptr.	Term ptr.
33		
29		
44		
126		
7		

} Save 9 bytes  
on 3  
pointers.

← Lose 4 bytes on  
term lengths.



# Net Result

- Example for block size  $k = 4$
- Where we used 3 bytes/pointer without blocking
  - $3 \times 4 = 12$  bytes,

now we use  $3 + 4 = 7$  bytes.


Shaved another  $\sim 0.5$  MB. This reduces the size of the dictionary from 7.6 MB to 7.1 MB.

We can save more with larger  $k$ .


Why not go with larger  $k$ ?

## Dictionary search without blocking

- Assuming each dictionary term equally likely in query (not true in practice!), average number of comparisons =  $(1 + (2 \cdot 2) + (4 \cdot 3) + 4) / 8$   
=  $\sim 2.6$



# Dictionary search with blocking



- Binary search down to 4-term block;
  - Then linear search through terms in block.
- Blocks of 4 (binary tree), average =  
 $(1+(2 \cdot 2)+(2 \cdot 3)+(2 \cdot 4)+5)/8 = 3$  compares



# Front coding

- Sorted words commonly have long common prefix – store differences only
  - Used in the (for last  $k-1$  in a block of  $k$ )

**8automata8automate9automatic10automation**

→ **8automat\*a1◇e2◇ic3◇ion**

Encodes **automat**

Extra length  
beyond **automat**.

*Begins to resemble general string compression*

# RCV1 dictionary compression summary

Technique	Size in MB
Fixed width	11.2
Dictionary-as-String with pointers to every term	7.6
Also, blocking $k = 4$	7.1
Also, Blocking + front coding	5.9



# POSTINGS COMPRESSION



# Postings compression

- The postings file is much larger than the dictionary, factor of at least 10.
- Key desideratum: store each posting compactly.
- A posting for our purposes is a **docID**.
- For Reuters (800,000 documents), we would use 32 bits per docID when using 4-byte integers.
- Alternatively, we can use  $\log_2 800,000 \approx 20$  bits per docID.
- Our goal: use a lot less than 20 bits per docID.





# Postings: two conflicting forces

- A term like ***arachnocentric*** occurs in maybe one doc out of a million – we would like to store this posting using  $\log_2 1M \sim 20$  bits.
- A term like ***the*** occurs in virtually every doc, so 20 bits/posting is too expensive.
  - Prefer 0/1 bitmap vector in this case



# Postings file entry

- We store the list of docs containing a term in increasing order of docID.
  - **computer**: 33,47,154,159,202 ...
- Consequence: it suffices to store *gaps*.
  - 33,14,107,5,43 ...
- Hope: most gaps can be encoded/stored with far fewer than 20 bits.

# Three postings entries



	encoding	postings list				
THE	docIDs	...	283042	283043	283044	283045 ...
	gaps		1	1	1	...
COMPUTER	docIDs	...	283047	283154	283159	283202 ...
	gaps		107	5	43	...
ARACHNOCENTRIC	docIDs	252000	500100			
	gaps	252000	248100			

# Variable length encoding



- Aim:
  - For *arachnocentric*, we will use  $\sim 20$  bits/gap entry.
  - For *the*, we will use  $\sim 1$  bit/gap entry.
- If the average gap for a term is  $G$ , we want to use  $\sim \log_2 G$  bits/gap entry.
- Key challenge: encode every integer (gap) with about as few bits as needed for that integer.
- This requires *variable length encoding*
- Variable length codes achieve this by using short codes for small numbers

# Variable Byte (VB) codes



- For a gap value  $G$ , we want to use close to the fewest bytes needed to hold  $\log_2 G$  bits
- Begin with one byte to store  $G$  and dedicate 1 bit in it to be a continuation bit  $c$
- If  $G \leq 127$ , binary-encode it in the 7 available bits and set  $c = 1$
- Else encode  $G$ 's lower-order 7 bits and then use additional bytes to encode the higher order bits using the same algorithm
- At the end set the continuation bit of the last byte to 1 ( $c = 1$ ) – and for the other bytes  $c = 0$ .

# Example

docIDs	824	829	215406
gaps		5	214577
VB code	00000110 10111000	10000101	00001101 00001100 10110001

$$512 + 256 + 32 + 16 + 8 = 824$$

Postings stored as the byte concatenation

00000110 10111000 10000101 00001101 00001100 10110001

Key property: VB-encoded postings are uniquely prefix-decodable.

For a small gap (5), VB uses a whole byte.

# Other variable unit codes



- Instead of bytes, we can also use a different “unit of alignment”: 32 bits (words), 16 bits, 4 bits (nibbles).
- Variable byte alignment wastes space if you have many small gaps – nibbles do better in such cases.
- Variable byte codes:
  - Used by many commercial/research systems
  - Good low-tech blend of variable-length coding and sensitivity to computer memory alignment matches (vs. bit-level codes, which we look at next).



# RCV1 compression

Data structure	Size in MB
dictionary, fixed-width	11.2
dictionary, term pointers into string	7.6
with blocking, $k = 4$	7.1
with blocking & front coding	5.9
collection (text, xml markup etc)	3,600.0
collection (text)	960.0
Term-doc incidence matrix	40,000.0
postings, uncompressed (32-bit words)	400.0
postings, uncompressed (20 bits)	250.0
postings, variable byte encoded	116.0





# Summary: Index compression

- We can now create an index for highly efficient Boolean retrieval that is very space efficient
- Use the sorted nature of the data to compress
  - Variable sized storage
  - Encode common prefixes only once
  - Encode gaps to reduce size of numbers
- However, here we didn't encode positional information
  - But techniques for dealing with postings are similar



# Resources for today's lecture

- *IIR 5*
- *MG 3.3, 3.4.*
- F. Scholer, H.E. Williams and J. Zobel. 2002. Compression of Inverted Indexes For Fast Query Evaluation. *Proc. ACM-SIGIR 2002.*
  - Variable byte codes
- V. N. Anh and A. Moffat. 2005. Inverted Index Compression Using Word-Aligned Binary Codes. *Information Retrieval 8*: 151–166.
  - Word aligned codes