

Information Retrieval and Organisation

Dell Zhang


Birkbeck, University of London

2015/16

Dictionaries and Tolerant Retrieval

Dictionaries

- Dictionary: the data structure for storing the term vocabulary



Storing Dictionaries

- ▶ For each term, we need to store a couple of items:
 - ▶ document frequency
 - ▶ pointer to postings list
 - ▶ ...
- ▶ Assume for the time being that
 - ▶ we can store this information in a fixed-length entry
 - ▶ we store these entries in an array

Storing Dictionaries

term	document frequency	pointer to postings list
a	656,265	→
aachen	65	→
...
zulu	221	→

space needed: 20 bytes 4 bytes 4 bytes

- ▶ How do we look up an element in this array at query time?
- ▶ Remember: these dictionaries can be huge, scanning is not an option

Data Structures

- ▶ Two main classes of data structures:
hash tables and trees
 - ▶ Some IR systems use hash tables, some use trees.
- ▶ Criteria for when to use hash tables vs trees:
 - ▶ Is there a fixed number of terms or will it keep growing?
 - ▶ What are the relative frequencies with which various keys will be accessed?
 - ▶ How many terms are we likely to have?

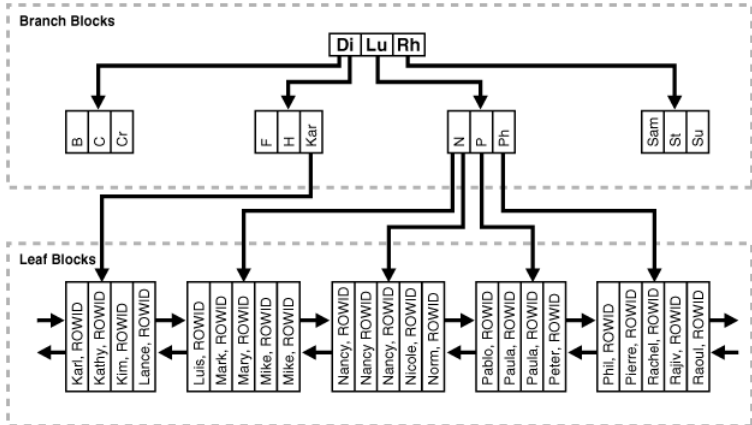
Hash Tables

- ▶ Each vocabulary term is hashed into an integer.
- ▶ Try to avoid collisions
- ▶ At query time, do the following:
 - ▶ hash query term
 - ▶ resolve collisions
 - ▶ locate entry in fixed-width array
- ▶ Pros:
 - ▶ Lookup in a hash table is faster than in a tree.
- ▶ Cons:
 - ▶ no prefix search (all terms starting with *automat*)
 - ▶ need to rehash everything periodically if vocabulary keeps growing

Trees

- ▶ Trees solve the prefix problem (find all terms starting with *automat*).
- ▶ Simplest tree: binary tree.
- ▶ However, binary trees are problematic:
 - ▶ Only balanced trees allow efficient retrieval
 - ▶ Rebalancing binary trees is expensive
- ▶ Use B-trees (the index structure that you know from database lectures)

B-Tree



Taken from documentation for Oracle 10g

Wildcard Queries

- ▶ `mon*`: find all docs containing any term beginning with *mon*
 - ▶ Easy with B-tree dictionary
 - ▶ retrieve all terms t in the range: $\text{mon} \leq t < \text{moo}$
- ▶ `*mon`: find all docs containing any term ending with *mon*
 - ▶ Maintain an additional tree for terms *backwards*, then
 - ▶ retrieve all terms t in the range: $\text{nom} \leq t < \text{non}$

Query Processing


- ▶ At this point, we have an enumeration of all terms in the dictionary that match the wildcard query.
- ▶ We still have to look up the postings for each enumerated term.
 - ▶ e.g., consider the query: `gen* AND universit*`
- ▶ This may result in the execution of many Boolean AND queries.

Wildcards in Middle of Term

- ▶ Example: `m*nchen`
- ▶ We could look up `m*` and `*nchen` in the B-tree and intersect the two term sets.
 - ▶ Expensive (there are probably thousands and thousands of terms beginning with “m”)
- ▶ Alternative: *permuterm* index
 - ▶ Basic idea: Rotate every wildcard query, so that the `*` occurs at the end.

Permuterm Index

- For term `hello`: add *hello\$*, *ello\$h*, *llo\$he*, *lo\$hel*, *o\$hell*, and *\$hello* to the B-tree where `$` is a special symbol



Permuterm Index

- ▶ Queries
 - ▶ For X , look up $X\$$
 - ▶ For X^* , look up $\$X^*$
 - ▶ For $*X$, look up $X\*
 - ▶ For $*X^*$, look up X^*
 - ▶ For X^*Y , look up $Y\$X^*$
- ▶ Example:
 - ▶ For hel^*o , look up $o\$hel^*$
- ▶ It's really a tree and should be called permuterm tree
- ▶ But permuterm index is more common name.


Query Processing

- ▶ Once we modified the query (as shown on last slide), we can do a regular lookup on a B-tree
- ▶ This is much faster than looking up X^* and $*Y$ and combining results (for query X^*Y)
- ▶ Permuterm index also handles leading wildcards: $*X$
- ▶ It has a disadvantage, though: quadruples the size of the dictionary compared to a regular B-tree (as every term is stored multiple times)

k -gram Index

- ▶ More space-efficient than permuterm index
- ▶ Enumerate all character k -grams (sequence of k characters) occurring in a term
 - ▶ 2-grams are also called *bigrams*
 - ▶ 3-grams are also called *trigrams*
- ▶ Example:
 - ▶ from *April is the cruelest month*
we get the bigrams:
 $\$a\ ap\ pr\ ri\ il\ l\$ \$i\ is\ s\$ \$t\ th\ he\ e\$ \$c\ cr\ ru\ ue\ el\ le$
 $es\ st\ t\$ \$m\ mo\ on\ nt\ h\$$
 - ▶ $\$$ is a special word boundary symbol.
- ▶ Maintain an inverted index from bigrams to the terms that contain the bigram

Postings List in a 3-gram Index



- ▶ Note that we now have two different types of inverted indexes
 - ▶ The term-document inverted index for finding documents based on a query consisting of terms
 - ▶ The k -gram index for finding terms based on a query consisting of k -grams

Processing Wildcard Queries

- ▶ Query `mon*` can now be run as:
`$m AND mo AND on`
- ▶ Gets us all terms with the prefix *mon* . . .
- ▶ . . . but also many “false positives” like `moon`
- ▶ We must post-filter these terms against query
- ▶ Surviving terms are then looked up in the term-document inverted index.
- ▶ *k*-gram indexes are fast and space efficient (compared to permuterm indexes).

Processing Wildcard Queries

- ▶ We must potentially execute a large number of Boolean queries for each enumerated, filtered term (on the term-document index)
 - ▶ Recall the query: `gen* AND universit*`
 - ▶ Most straightforward semantics: Conjunction of disjunctions
 - ▶ Very expensive
- ▶ Users hate to type
 - ▶ If abbreviated queries like `pyth* theo*` for `pythagoras' theorem` are legal, users will use them ...
 - ▶ ... a lot

Spelling Correction

- ▶ Two principal uses
 - ▶ Correcting documents being indexed
 - ▶ Correcting user queries
- ▶ Two different methods
 - ▶ *Isolated Word* Spelling Correction
 - ▶ Check each word on its own for misspelling
 - ▶ Will not catch typos resulting in correctly spelled words, e.g., *an asteroid that fell form the sky*
 - ▶ *Context-Sensitive* Spelling Correction
 - ▶ Look at surrounding words
 - ▶ Can correct the *form/from* error above

Correcting Documents

- ▶ We're not interested in interactive spelling correction of documents (e.g., MS Word) in this class.
- ▶ In IR, we use document correction primarily for OCR'ed documents (i.e. documents digitized via Optical Character Recognition)
- ▶ The general philosophy in IR is: don't change the documents.

Correcting Queries

- ▶ First: isolated word spelling correction
 - ▶ Fundamental premise 1: There is a list of “correct words” from which the correct spellings come.
 - ▶ Fundamental premise 2: We have a way of computing the *distance* between a misspelled word and a correct word.
- ▶ Simple spelling correction algorithm:
return the “correct” word that has the *smallest* distance to the misspelled word.
 - ▶ Example: *informaton* → *information*

Correcting Queries

- ▶ Can we use the term vocabulary of the inverted index as the list of correct words?
 - ▶ It can be very biased
 - ▶ It may be missing certain terms
- ▶ Alternatives:
 - ▶ A standard dictionary
(Webster's, Encyclopædia Britannica, etc.)
 - ▶ An industry-specific dictionary
(for specialized IR systems)
 - ▶ The term vocabulary of the collection,
appropriately weighted

Computing Distance

- ▶ How can we compute the distance between words?
- ▶ We'll look at some alternatives:
 - ▶ edit distance (Levenshtein distance)
 - ▶ weighted edit distance
 - ▶ k -gram overlap

Edit Distance

- ▶ The (minimum) edit distance between two strings s_1 and s_2 is the minimum number of basic operations to convert s_1 to s_2 .
- ▶ Levenshtein distance: the admissible basic operations are: insert, delete, and replace
 - ▶ Levenshtein distance $dog \rightarrow do$: 1 (deletion)
 - ▶ Levenshtein distance $cat \rightarrow cart$: 1 (insertion)
 - ▶ Levenshtein distance $cat \rightarrow cut$: 1 (replacement)
 - ▶ Levenshtein distance $cat \rightarrow act$: 2
(2 replacements or 1 insertion and 1 deletion)

Computing Distance

- ▶ Getting from *cats* to *fast*

	""	f	a	s	t
""	"" → ""	"" → f	"" → fa	"" → fas	"" → fast
c	c → ""	c → f	c → fa	c → fas	c → fast
a	ca → ""	ca → f	ca → fa	ca → fas	ca → fast
t	cat → ""	cat → f	cat → fa	cat → fas	cat → fast
s	cats → ""	cats → f	cats → fa	cats → fas	cats → fast


- ▶ Each cell will contain the (cheapest) cost of getting from the string on the left-hand side to the string on the right-hand side

Computing Distance

- ▶ We know the costs for the uppermost row and the leftmost column:
 - ▶ we have to get from "" to *fast* by inserting characters
 - ▶ we have to get from *cats* to "" by deleting characters

	""	f	a	s	t
""	0	1	2	3	4
c	1				
a	2				
t	3				
s	4				

Computing Distance



- ▶ For other cells, take the minimum of costs
 - ▶ Coming from (a):
 - ▶ add 1 to cost in (a) — insertion
 - ▶ Coming from (b):
 - ▶ add 1 to cost in (b) — deletion
 - ▶ Coming from (c):
 - ▶ if characters in row and column are equal, copy cost from (c)
 - ▶ otherwise, add 1 to cost in (c) — replacement

Resulting Matrix

- ▶ Computing the costs for all cells results in the following matrix:

	''	f	a	s	t
''	0	1	2	3	4
c	1	1	2	3	4
a	2	2	1	2	3
t	3	3	2	2	2
s	4	4	3	2	3

- ▶ So the Levenshtein distance is 3

Algorithm

EDITDISTANCE(s_1, s_2)

```
1  int  $m[i, j] = 0$ 
2  for  $i \leftarrow 1$  to  $|s_1|$ 
3  do  $m[i, 0] = i$ 
4  for  $j \leftarrow 1$  to  $|s_2|$ 
5  do  $m[0, j] = j$ 
6  for  $i \leftarrow 1$  to  $|s_1|$ 
7  do for  $j \leftarrow 1$  to  $|s_2|$ 
8      do  $m[i, j] = \min\{m[i-1, j-1] + \text{if } (s_1[i] = s_2[j]) \text{ then } 0 \text{ else } 1, \text{fi},$ 
9           $m[i-1, j] + 1,$ 
10          $m[i, j-1] + 1\}$ 
11 return  $m[|s_1|, |s_2|]$ 
```

Weighted Edit Distance

- ▶ As Levenshtein distance, but weight of an operation depends on the characters involved.
- ▶ Meant to capture keyboard errors
 - ▶ e.g., m more likely to be mistyped as n than as q .
 - ▶ therefore, replacing m by n is a smaller edit distance than by q .
- ▶ We now require a weight matrix as input.
- ▶ Modify dynamic programming to handle weights.

Using Edit Distances

- ▶ Comparing query term q to all terms in the vocabulary is too expensive
- ▶ Solution: use heuristics to determine subset
 - ▶ Only compare to terms beginning with the same letter (doesn't work for typos at beginning)
 - ▶ Generate set of rotations for q and use a permuterm index (doesn't work well for replacements)
 - ▶ For each rotation, omit a suffix of l characters before doing lookup in permuterm index
 - ▶ Ensures that each term in query rotation shares a substring with retrieved terms
 - ▶ The value of l could be fixed to a constant length (e.g. 2), or depend on the length of q


Using a k -gram Index

- ▶ Enumerate all k -grams in the query term
- ▶ Use the k -gram index to retrieve “correct” words that match query term k -grams
- ▶ Threshold by number of matching k -grams
 - ▶ e.g., only vocabulary terms that differ by at most 3 k -grams

Example with 2-grams

- Suppose the misspelled word is “bordroom”:

\$b, bo, or, rd, dr, ro, oo, om, m\$



Example with 3-grams

- ▶ Suppose the correct word is “november”:
\$\$n, \$no, nov, ove, vem, emb, mbe, ber, er\$, r\$\$
- ▶ And the query term is “december”:
\$\$d, \$de, dec, ece, cem, emb, mbe, ber, er\$, r\$\$
- ▶ So 5 trigrams overlap (out of 10 in each term)
- ▶ Issue: Fixed number of k -grams that differ does not work for words of differing length.
- ▶ How can we turn this into a normalized measure of overlap?

Jaccard Coefficient

- ▶ A commonly used measure of two sets' overlap
- ▶ Let A and B be two sets
- ▶ Jaccard coefficient:

$$\frac{|A \cap B|}{|A \cup B|}$$

- ▶ A and B don't have to be the same size.
 - ▶ Always assigns a number between 0 and 1.
- ▶ Application to spelling correction: declare a match if the coefficient is, say, > 0.8 .

Context-Sensitive Correction

- ▶ Our example was:
“an asteroid that fell *form* the sky”
- ▶ How can we correct *form* here?
- ▶ One idea: *hit-based* spelling correction
 - ▶ We'll return back to this idea when we talk about the *probabilistic* approach to spelling correction, in the second half of the module.

Context-Sensitive Correction

- ▶ Given query “flew *form* munich”
- ▶ Retrieve the correct terms close to each query term
 - ▶ flea for *flew*
 - ▶ from for *form*
 - ▶ munch for *munich*
- ▶ Now try all possible resulting phrases as queries, with one word fixed at a time
 - ▶ Try query “flea form munich”
 - ▶ Try query “flew from munich”
 - ▶ Try query “flew form munch”
- ▶ The correct query “flew from munich” should have the most hits.

Context-Sensitive Correction

- ▶ The *hit-based* algorithm we just outlined is not very efficient.
 - ▶ Suppose we have 7 alternatives for *flew*, 19 for *form* and 3 for *munich*
 - ▶ Then we have to test $7 \times 19 \times 3$ different variants
- ▶ More efficient alternative: look at the collection of queries, not documents
 - ▶ This assumes that we log queries

General Issues

- ▶ User interface
 - ▶ Automatic or suggested correction
 - ▶ “*Did you mean*” only works for one suggestion.
 - ▶ What about multiple possible corrections?
 - ▶ Tradeoff: simple vs powerful UI
- ▶ Cost
 - ▶ Spelling correction is potentially expensive.
 - ▶ Avoid running on every query?
 - ▶ Maybe just on queries that match few documents.

Phonetic Matching

- ▶ Soundex is the basis for finding *phonetic* (as opposed to orthographic) alternatives.
 - ▶ e.g., Chebyshev / Tchebyscheff
- ▶ Algorithm:
 - ▶ Turn every token to be indexed into a 4-character reduced form
 - ▶ Do the same with query terms
 - ▶ Build and search an index on the reduced forms

Soundex Algorithm

1. Retain the first letter of the term.
2. Change all occurrences of the following letters to 0 (zero):
 - ▶ A, E, I, O, U, H, W, Y
3. Change letters to digits as follows:
 - ▶ B, F, P, V \Rightarrow 1
 - ▶ C, G, J, K, Q, S, X, Z \Rightarrow 2
 - ▶ D, T \Rightarrow 3
 - ▶ L \Rightarrow 4
 - ▶ M, N \Rightarrow 5
 - ▶ R \Rightarrow 6
4. Repeatedly remove one out of each pair of consecutive identical digits
5. Remove all 0s from the resulting string; pad the resulting string with trailing 0s, and return the first four positions, which will consist of a letter followed by three digits

Soundex Algorithm

- ▶ Example

	difficulty	difference
steps 1 and 2	d0ff0c0lt0	d0ff0r0nc0
step 3	d011020430	d011060520
step 4	d01020430	d01060520
step 5	d124	d165

- ▶ Vowels are viewed as being interchangeable
- ▶ Consonants with similar sounds (e.g. D and T) are put in equivalence classes
- ▶ Works fairly well for European languages

Summary

- ▶ How to organize a dictionary of an inverted index
- ▶ How to do imprecise searches on this dictionary handling
 - ▶ wildcards
 - ▶ spelling mistakes